

# parallel programming project – 2d incompressible fluid


Nicole Vadot

# initial serial code

$$\nabla \cdot v = 0$$

$$\frac{\partial \mathbf{u}}{\partial t} = -(\mathbf{u} \cdot \nabla) \mathbf{u} + \nu \nabla^2 \mathbf{u} + \mathbf{f}$$
$$\frac{\partial \rho}{\partial t} = -(\mathbf{u} \cdot \nabla) \rho + \kappa \nabla^2 \rho + S$$

- diffusion → Gauss-Seidel relaxation
- advection → particle method
- incompressibility → "projection" (also GS)
- complexity  $\sim nx \cdot ny \cdot nrelax$
- $nrelax = cst$  to prevent implicit dependence on  $nx, ny$  (no threshold)



```
void dens_step ( uint nx, uint ny, uint nrelax, double * x, double * x0, double * u,
    add_source ( nx, ny, x, x0, dt );
    SWAP ( x0, x ); diffuse ( nx, ny, nrelax, 0, x, x0, diff, dt );
    SWAP ( x0, x ); advect ( nx, ny, 0, x, x0, u, v, dt );
}

void vel_step ( uint nx, uint ny, uint nrelax, double * u, double * v, double * u0, d
    add_source ( nx, ny, u, u0, dt ); add_source ( nx, ny, v, v0, dt );
    SWAP ( u0, u ); diffuse ( nx, ny, nrelax, 1, u, u0, visc, dt );
    SWAP ( v0, v ); diffuse ( nx, ny, nrelax, 2, v, v0, visc, dt );
    project ( nx, ny, nrelax, u, v, u0, v0 );
    SWAP ( u0, u ); SWAP ( v0, v );
    advect ( nx, ny, 1, u, u0, u0, v0, dt ); advect ( nx, ny, 2, v, v0, u0, v0, dt );
    project ( nx, ny, nrelax, u, v, u0, v0 );
}
```

```
typedef struct {
    uint nx, ny;
    double diff, visc, dt;
    uint nrelax; // TODO : use a threshold instead if this is -1 ?
    double *rho, *rho_prev, *vx, *vx_prev, *vy, *vy_prev;
} Sim;
```

# where to optimize ?

```
make -B USE_OMP=0 USE_GS=0 RELEASE=1 main_perf && perf record ./main_perf 100 100
```

```
Samples: 21K of event 'cycles:Pu', Event count (approx.): 23797168187
```

Overhead	Command	Shared Object	Symbol
91,42%	main_perf	main_perf	[.] lin_solve
1,39%	main_perf	main_perf	[.] advect
1,16%	main_perf	main_perf	[.] project
0,84%	main_perf	main_perf	[.] set_bnd
0,58%	main_perf	libm.so.6	[.] 0x00000000000075f0a
0,56%	main_perf	libm.so.6	[.] 0x00000000000073c84
0,52%	main_perf	main_perf	[.] main

```
void lin_solve (uint nx, uint ny, uint nrelax, int b, double *  
x, const double * x0, double a, double c ) {  
    for (uint k=0 ; k<nrelax ; ++k) {  
        for (uint i=1 ; i<nx-1 ; ++i) {  
            for (uint j=1 ; j<ny-1 ; ++j) {  
                x[IX(i,j)] = (x0[IX(i,j)] + a*(x[IX(i-1,j)]+x[IX(i+1,j)]  
                    +x[IX(i,j-1)]+x[IX(i,j+1)]))/c;  
            }  
        }  
        set_bnd (nx, ny, b, x);  
    }  
}
```

⇒ most of the time is spent running lin\_solve

# 1. optimize lin\_solve

works: `#pragma omp parallel for`

slower ??? : `#pragma omp parallel for collapse(2)`

~similar: `#pragma omp parallel for schedule(static, (nx-2)/num_threads))`

```
void lin_solve (uint nx, uint ny, uint nrelax, int b, double * x, con
#ifdef _OPENMP
int num_threads;
#pragma omp parallel
{ num_threads = omp_get_num_threads(); }
#endif
for (uint k=0 ; k<nrelax ; ++k) {
#ifdef _OPENMP
#pragma omp parallel for schedule(static, (nx-2)/num_threads)
#endif
for (uint i=1 ; i<nx-1 ; ++i) {
for (uint j=1 ; j<ny-1 ; ++j) {
x[IX(i,j)] = (x0[IX(i,j)] + a*(x[IX(i-1,j)]+x[IX(i+1,j)]+x[IX
}
}
set_bnd (nx, ny, b, x);
}
```

`nx=100, ny=100, nrelax=100`

no omp: ~30ms/step

1 thread: ~32ms/step overhead!

2 threads: ~16ms/step

3 threads: ~12ms/step

4 threads: ~9.2ms/step

5 threads: ~7.9ms/step

6 threads: ~6.7ms/step

7 threads: ~5.8ms/step strong scaling

8 threads: ~5.9ms/step limit!

9 threads: ~7.3ms/step

...

# fixing the race condition


- problem: results are not repeatable !!
- race condition between the threads

```
for (uint k=0 ; k<nrelax ; ++k) {  
    #pragma omp parallel for schedule(static, (nx-2)/num_threads)  
    for (uint i=1 ; i<nx-1 ; ++i) {  
        for (uint j=1 ; j<ny-1 ; ++j) {  
            x[IX(i,j)] = (x0[IX(i,j)] + a*(x[IX(i-1,j)]+x[IX(i+1,j)]  
                +x[IX(i,j-1)]+x[IX(i,j+1)]))/c;  
        }  
    }  
    set_bnd (nx, ny, b, x);  
}
```


- use Jacobi method instead of Gauss-Seidel

```
for (uint k=0 ; k<nrelax ; ++k) {  
    #pragma omp parallel for schedule(static, (nx-2)*(ny-2)/num_threads)  
    for (uint i=1 ; i<nx-1 ; ++i) {  
        for (uint j=1 ; j<ny-1 ; ++j) {  
            x_[IX(i,j)] = (x0[IX(i,j)] + a*(x[IX(i-1,j)]  
                +x[IX(i+1,j)]+x[IX(i,j-1)]+x[IX(i,j+1)]))/c;  
        }  
    }  
    SWAP(x, x_);  
    set_bnd(nx, ny, b, x);  
}
```

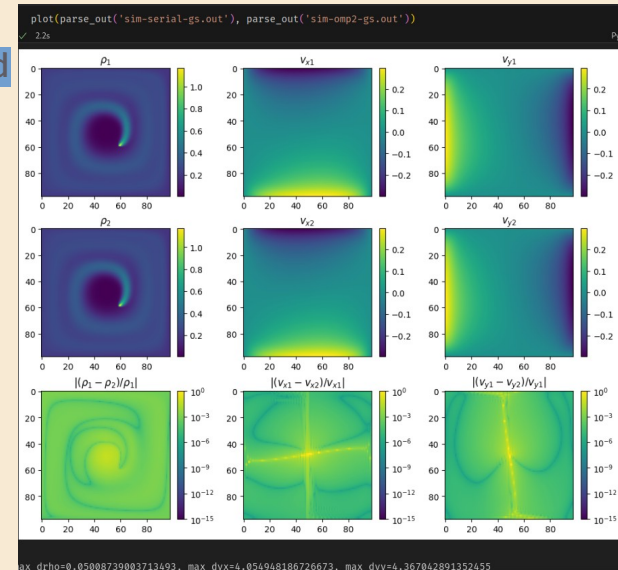
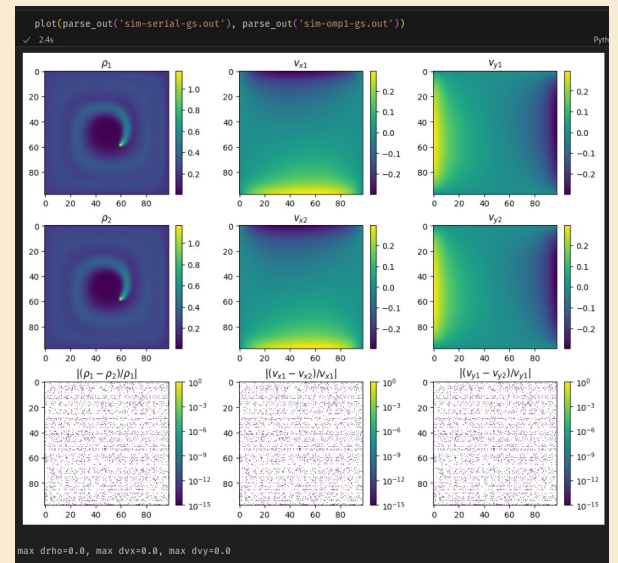
normally:



threads overlap!



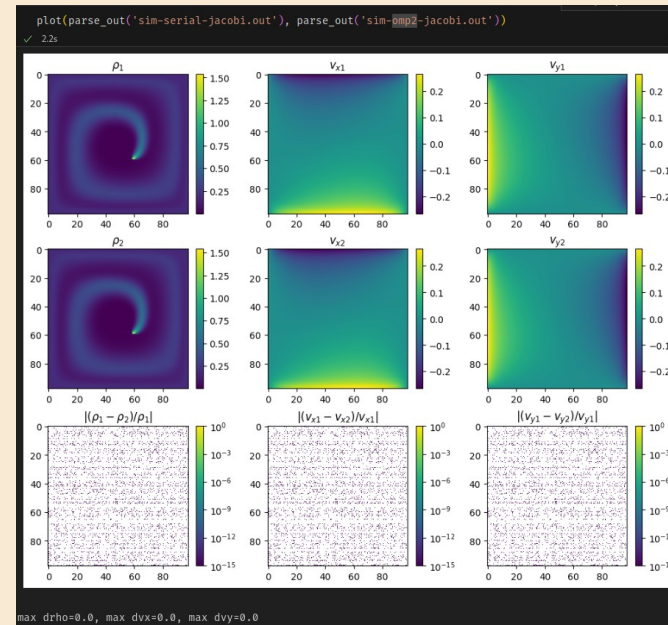
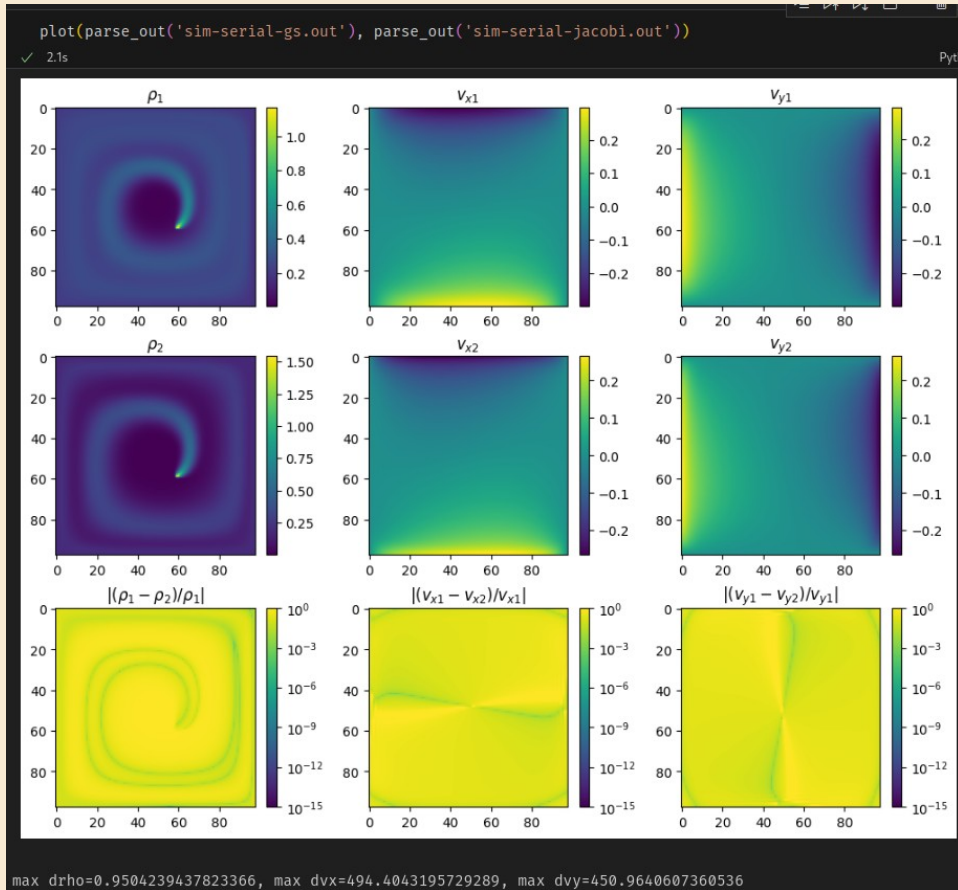
thread boundary



# jacobi vs gauss-seidel

→ at fixed nrelax, not the same!  
(different convergence rate)

→ jacobi ~5.3ms/step  
gs ~30ms/step  
→ ~6x speedup  
→ no more read-after-write!



timings with  
jacobi  
[ms/step]

noomp	5.3
1 th.	5.4
2 th.	3.4
3 th.	2.6
4 th.	2.2
5 th.	2.0
6 th.	1.9
7 th.	2.8




## 2. improve balance between density and velocity computation

- 5.3ms/step = 1.0ms/step [density]  
+ 4.3ms/step [velocity]
- serial algo (leapfrog-style):  $\rho \rightarrow v_x, v_y \rightarrow \rho \rightarrow v_x, v_y \rightarrow \dots$
- parallel algo:  $(\rho, v_x, v_y) \rightarrow (\rho, v_x, v_y) \rightarrow \dots$

new serial implementation vvv

```
void sim_step(Sim* sim) {  
    // make copy of vx, vy so that density can work independantly  
    memcpy(sim->vx_, sim->vx, sim->nx*sim->ny*sizeof(*sim->vx));  
    memcpy(sim->vy_, sim->vy, sim->nx*sim->ny*sizeof(*sim->vy));  
  
    // parallel region 1  
    vel_step(sim->nx, sim->ny, sim->nrelax, sim->vx, sim->vy, sim->vx_prev, sim->vy_prev, sim->visc, sim->dt);  
  
    // parallel region 2  
    dens_step(sim->nx, sim->ny, sim->nrelax, sim->rho, sim->rho_prev, sim->vx_, sim->vy_, sim->diff, sim->dt);  
}
```

compare with before.. →  
difference is  $\sim 1e-7$   
→ seems OK



# MPI parallel implementation

2 processes: dens  
and vel, each  
handles its own  
evolution

communications are  
only one-  
directional!  
(vel is indep of  
dens)

```
MPI_Request request_v[2];

long long t0 = current_timestamp();
long long totaltime = 0;

for (uint istep = 0; istep < nsteps; ++istep) {
    if (istep % 100 == 0) printf("[sim %d/%d] iter %u/%u\n", rank, size, istep, nsteps);

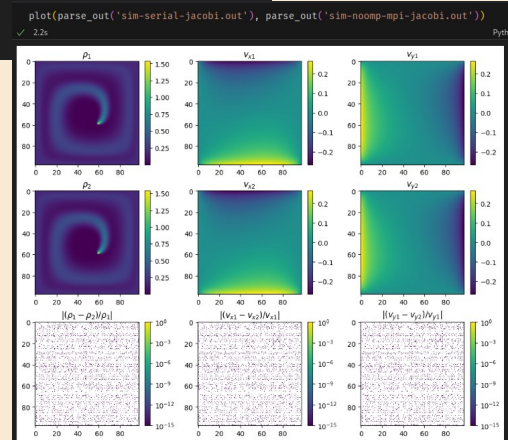
    long long timestep0 = current_timestamp();
    // apply sources and step
    if (role == ROLE_DENS) {
        spinny_dens(sim);
        dens_step(sim->nx, sim->ny, sim->nrelax, sim->rho, sim->rho_prev, sim->vx, sim->vy, sim->diff, sim->dt);
    } else {
        spinny_vel(sim);
        vel_step(sim->nx, sim->ny, sim->nrelax, sim->vx, sim->vy, sim->vx_prev, sim->vy_prev, sim->visc, sim->dt);
    }
    totaltime = current_timestamp() - timestep0;

    // vel process needs to update the vx,vy buffers of the density process
    // but vel is indep of density
    if (role == ROLE_DENS) {
        MPI_Irecv(sim->vx, sim->nx*sim->ny, MPI_DOUBLE, rank_of_other, 0, MPI_COMM_WORLD, &request_v[0]);
        MPI_Irecv(sim->vy, sim->nx*sim->ny, MPI_DOUBLE, rank_of_other, 0, MPI_COMM_WORLD, &request_v[1]);
    } else {
        MPI_Isend(sim->vx, sim->nx*sim->ny, MPI_DOUBLE, rank_of_other, 0, MPI_COMM_WORLD, &request_v[0]);
        MPI_Isend(sim->vy, sim->nx*sim->ny, MPI_DOUBLE, rank_of_other, 0, MPI_COMM_WORLD, &request_v[1]);
    }
    MPI_Waitall(2, request_v, MPI_STATUS_IGNORE);
}
```

verify identical results [OK] →

$$\frac{\partial \mathbf{u}}{\partial t} = -(\mathbf{u} \cdot \nabla) \mathbf{u} + \nu \nabla^2 \mathbf{u} + \mathbf{f}$$

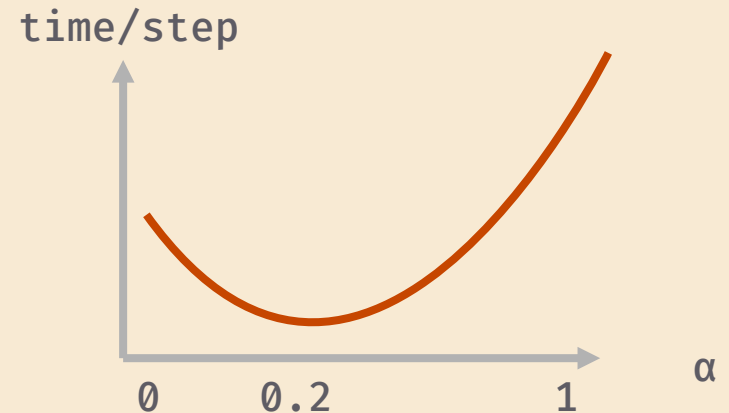
$$\frac{\partial \rho}{\partial t} = -(\mathbf{u} \cdot \nabla) \rho + \kappa \nabla^2 \rho + S$$





# balancing the dens, vel MPI processes (the theory)

- $5.3\text{ms/step} = 1.0\text{ms/step} [\text{density}] + 4.3\text{ms/step} [\text{velocity}]$   
⇒ time/step is limited by the longest of dens, vel process
  - idea: give more threads to velocity than density to equilibrate?
  - $n_t$  threads available,  $0 < \alpha < 1$ 
    - $\alpha * n_t$  go to density
    - $(1-\alpha) * n_t$  go to velocity(+/- some rounding and thresholding)
  - what  $\alpha$  to pick?
    - do a scan to minimize time/step
    - intuition -- 4:1 ratio
      - velocity: 4 calls to lin\_solve per step
      - density: 1 call to lin\_solve per step
- ⇒  $\alpha \sim 1/(4+1) = 0.2$



# balancing the dens, vel MPI processes (the practice)

- didn't get it to work :(  
enabling OMP on top of MPI results in slower times (and more threads leads to even slower)
- tried tweaking
  - `MPI_THREAD_LEVEL, MPI_Init{,_thread}`
  - OMP parallel for collapse, schedule ...
  - number of threads, grid size ...

## further improvements

- first touch: OK probably
- MPI cartesian geometry  
...but, might help on large grids (not the case here, we are 2D)
- prevent fork/join on every lin\_solve  
+ would require to also parallelize all other loops  
(or add omp single everywhere else)
- review vectorized instructions  
...seems OK... maybe cache misses could be reduced?
- interface parallelized MPI with GUI (once MPI+OMP works)

```
Sim* sim_new(uint nx, uint ny, double dt, double diff, double visc, uint nrelax) {
    Sim* sim = malloc(sizeof(*sim));
    sim->nx = nx;
    sim->ny = ny;
    sim->dt = dt;
    sim->diff = diff;
    sim->visc = visc;
    sim->nrelax = nrelax;
    sim->vx = calloc(nx*ny, sizeof(*sim->vx));
    sim->vx_prev = calloc(nx*ny, sizeof(*sim->vx_prev));
    sim->vy = calloc(nx*ny, sizeof(*sim->vy));
    sim->vy_prev = calloc(nx*ny, sizeof(*sim->vy_prev));
    sim->rho = calloc(nx*ny, sizeof(*sim->rho));
    sim->rho_prev = calloc(nx*ny, sizeof(*sim->rho_prev));
    return sim;
}
```

[illegible]

```
objdump -Intel --visualize-jumps --disassemble=lin_solve main_perf
```